

Before You Believe It

A five-to-seven-minute family activity that turns a confident AI answer into a shared practice of pausing, questioning, checking, and thinking together.

01

Think

The child chooses the question that makes a plausible claim worth pausing over.

02

Push Back

The family reveals the hidden assumption doing work inside the claim.

03

Check

They connect a clue or counter-perspective—not a verdict—to the claim.

04

Make

The child assembles a provisional answer in their own words.

05

Own

A Learning Receipt separates child work, AI help, evidence, uncertainty, and a parent question.

Static first.

The complete Floating City MissionPack works without a key or network.

AI is bounded.

GPT-5.6 prepares material to question; it never grades, ranks, diagnoses, or decides.

Child work stays local.

The live request receives only a general parent topic and age band.

Why it matters

Parents want to understand how children learn with AI—not only what answer they repeat. Before You Believe It creates an approachable weekend moment for that conversation.

BUILT AND VERIFIED WITH

Codex + GPT-5.6

Codex accelerated the accessible canvas, tests, deployment boundary, and evidence. GPT-5.6 powers the optional live challenge generator.

31 unit tests

Static fallback

before-you-believe-it-build.vercel.app

Pause · Question · Check · Think together

One bounded model call, inside a reliable family activity.

The preset mission always runs locally. The optional reviewer demo adds a server-side GPT-5.6 challenge without exposing credentials or a child’s answer.

FAMILY BROWSER

Reasoning canvas

- Static Floating City MissionPack
- Think → Push Back → Check → Make → Own
- Session-only child response and receipt

No API key. Child answer is never included in a model request.

ALWAYS AVAILABLE

Preset fallback

A complete mission remains usable if network, access code, or model availability fails.



VERCEL SERVERLESS BOUNDARY

Verify and validate

- Judge code → signed HTTP-only cookie
- Validate parent topic + age band
- Strict JSON Schema and output validation

Malformed or unsafe model output never reaches the canvas.

OPENAI RESPONSES API

GPT-5.6

- Creates claim, hidden assumption, three clues, parent prompt, uncertainty
- Low reasoning effort + structured output
- Forbidden: grading, scoring, diagnosis, ranking, persuasion, verdicts

Privacy boundary: child choices and working answer are browser-session only.

Safety boundary: schema validation constrains the browser contract.

Reliability boundary: static mission is not dependent on a paid model call.


Actual implementation: static HTML/CSS/JS + Vercel functions

Optional live GPT-5.6 reviewer demo

LIVE-REQUEST HANDSHAKE · ACTUAL IMPLEMENTATION

How one live challenge becomes safe material to question.

The human authorizes a general-topic request. The server validates both access and output. The model never sees the child’s working answer.



		returning bounded fields.	
⑥ Family questions the returned material and makes a child-owned answer.	⑤ Renders validated text nodes; preserves static fallback if unavailable.		

Only a parent’s general topic and age band cross the model boundary.


Static MissionPack remains complete without live access.

VALIDATED MISSIONPACK WORKFLOW

FUTURE ROADMAP · NOT THE CURRENT RUNTIME

A future GPT-5.6 manager flow, with a human-safe release gate.

The shipped demo uses one structured GPT-5.6 call. This proposed server-side Agents SDK topology expands content preparation without changing the browser’s safety boundary.



Evidence Guardian

Checks source-aware evidence cards and uncertainty bounds.

DETERMINISTIC RELEASE GATE

MissionPack validator

- Schema and age/reading-level checks
- Provenance and evidence requirements
- Reject scoring, diagnosis, persuasion, and answer judgment
- Apply human-reviewed editorial rules

Important distinction

This is the documented next step, not a claim about today’s runtime. The current browser already receives only schema-validated material and retains a bundled fallback.

Fan-out specialists → validation → static canvas renderer

Future server-side GPT-5.6 Agents SDK architecture